

QUERYPILOT LOCAL

Evidence first.
LLM second.

Offline PostgreSQL plan analysis with a
deterministic correctness path.



The fast path ends before the model starts

1. VALIDATE

2. PROVE

3. EXPLAIN

Read-only SQL

Plan evidence

Optional wording

AST gate + 3 s
database timeout

Normalize JSON +
apply deterministic
rules

Retrieve docs + select
approved sentences

Two runtimes share one trust boundary

LOCAL RUNTIME

PostgreSQL + FastAPI +
rules + optional Foundry.

PUBLIC DEMO

Browser-only plan JSON;
no database or model.

Input

Core

Optional

Public

One read-only SELECT

Parser + rule engine

Local RAG + sentence IDs

EXPLAIN JSON in browser

No evidence means no recommendation

Reliability is enforced in code

Category, severity, plan evidence, SQL and citations are application-owned.

Unknown IDs, invented citations, extra prose and malformed output fall back safely.

- ✓ Single read-only SELECT
- ✓ Plan evidence required
- ✓ Unknown citations rejected
- ✓ Recommendation SQL is display-only
- ✓ No-answer is a valid result

The MVP fixtures pass; CPU enrichment stays optional

MEASURED RESULT

12 committed scenarios. Citations were valid in 4/4 generation samples; the 1.5B selector was accepted in 4/4.

12/12

Rule diagnosis
and no-answer

9/9

Retrieval
Hit@3

20.2 s

Average optional
CPU selection

EVIDENCE

Proof measured. Limits explicit.

1.671 ms Seq Scan → 0.074 ms Index Scan (7-run medians)

Local fixture only · optional CPU enrichment

Next: tool-calling v2 · querypilot.eraykulkizaga.com

